

markets

One row per market:

- market_id (uuid)
- market_name, city, state

properties_core

Stable identity + location fields:

- property_id (pk)
- market_id (fk)
- title, city, state, zipcode
- latitude, longitude
- listing URLs, host URL if needed

performance

Time/financial performance fields:

- property_id (pk/fk)
- revenue, revenue_potential, adr, occupancy, etc.
- total_reviews, property_rating (important later for charts)

property_features

Amenity + attribute flags:

- property_id (pk/fk)
- bedrooms, accommodates, bathrooms
- has_aircon, has_pool, etc.
- “system_*” amenities

property_scores

Computed outputs:

- property_id (pk/fk)
- market_id
- investment_score (0–100)
- score_breakdown (JSON explainability)
- optionally computed_at

Why this split is good engineering

- performance can change, features change, core rarely changes
- scoring can be recomputed repeatedly without touching raw tables
- API can fetch “just what it needs” per endpoint